

GH repo: [https://github.com/evgerritz/cpsc429\\_labs/tree/1\\_p3](https://github.com/evgerritz/cpsc429_labs/tree/1_p3)

“make run”: Build, insert kernel module, run threads program with  $w=5$ ,  $n=150000$

## Part 3:

---

- Acknowledgements
  - Used this for information on pthreads, mutexes:  
<https://www.cs.cmu.edu/afs/cs/academic/class/15492-f07/www/pthreads.html>
- Overview
  - threads.c contains the code for running workers
    - Usage: `./threads w n`
  - `create_workers` uses `pthread_create` to start  $w$ , each of which run `do_work`
  - `do_work` is the function each thread begins executing
    - It does the following  $n$  times:
      - Get the current value of the counter
      - Set the value to the original value plus one
    - Later, I added pthread mutexes to synchronize reads and writes.
  - At the end the program checks to see if the value of the counter is the value that would be obtained if all instructions were performed sequentially
    - If it isn't, it outputs the percent error
- Challenges
  - Math in C is tricky. I tried to calculate percent error (as a double) from the actual and correct value for the counter (both `uint64_t`) using “double  $q = (\text{actual} - \text{expected}) / ((\text{double}) \text{expected});$ ” and this always resulted in a huge, obviously incorrect value. The solution was to cast the values to doubles and then do the calculation, though I am still not entirely sure why.
  - Another small challenge was in implementing the lock. Initially I did so without fully understanding the race condition, so I make sure all reads and write were individually atomic (i.e., the mutex locking and unlocking was local to the `get_` and `set_counter` functions). As spelled out in the discussion below, the race condition actually arises from not protecting a larger critical section.
- Questions
  - Document at what values of  $W$  you receive incorrect results. What is the exact reason the incorrect results occur? Be specific.
    - Using  $n = 150,000$ , I received incorrect results on all values of  $W > 1$ . The incorrect results occur because the threads do not always write (line 2 in below image) the incremented value of the counter immediately after reading its current value (line 1). Without, for example, mutexes, one thread may read the current value of  $x$ , after which the scheduler switches execution to another thread and possibly even more threads, ultimately changing the value to  $x + a$ . Once the scheduler returns to the thread with a variable `current_val`  $= x$ , it writes the

value  $x+1$  into memory, as it has no way of knowing that the actual value of the counter in memory (and not outdated the value of *current\_val*) is  $x + a$ .

```
void * do_work(void * n_ptr) {
    int n = *((int *) n_ptr);
    int i;
    uint64_t current_val, read_val;
    for (i=0; i<n; i++) {
        1. current_val = get_counter();
        2. set_counter(current_val+1);
    }
}
```

- 
- Describe your solution in the lab report and explain any pros/cons it has.
  - I used the pthread implementation of mutexes to ensure that no context switching could occur within an iteration of the read-write loop. This ensures that each reading from memory is immediately followed by writing the incremented value into memory. The read-write section becomes atomic, so only one thread can execute those instructions at once. This is exactly the behavior we wanted to have initially. The pros of this method are that it's simple, intuitive, and it works well. One con is that this implementation results in a new global variable (the alternative to which would be adding parameters to functions and passing tramp data). A more serious con is that there is no guarantee that the scheduler won't preempt this thread from within the critical section. We definitely don't want this to happen, as it will waste cycles executing other threads which will get stuck waiting for the lock until execution returns to the thread with access.

```
for (i=0; i<n; i++) {
    pthread_mutex_lock(&mymem_lock);
    current_val = get_counter();
    set_counter(current_val+1);
    pthread_mutex_unlock(&mymem_lock);
}
```

- 
- How could you avoid this problem from within kernel space? In other words, how could your memory driver module ensure that user programs get the correct result no matter how many threads are simultaneously accessing it?
  - Here is the issue: in C, arbitrary user code may call the memory access functions in an arbitrary order, with differing intentions for which sections should be executed atomically. Since the driver doesn't know which request comes from which thread, it needs to know which sections of instructions are expected to run in parallel. Simply adding a mutex to each access function would not work, as it


would not solve the issue of user code getting preempted within a critical section. We don't want a mutex to ensure only one function is executed at a time, but rather a mutex to ensure only one critical section is running at a time. Thus, one solution is to add an interface to the user for a locking mechanism in the kernel. Essentially, we are just moving the user-space mutex into kernel space and providing the user with some access. Another solution would be to write the module using Rust because then the compiler would guarantee that the module is thread safe. The simplest solution would probably be to add new access functions like, say, `alter(fd, f, len)` that could read `len` bytes and replace it with `f` performed on those bytes. One could then wrap this in a mutex to ensure no two kernel threads are executing this part of the driver at once. This method would allow users to perform (an admittedly specific set of) concurrent operations without worrying about mutexes.

- What other kinds of unsafe behavior does your module allow to happen? For example, what happens if you close the file descriptor before all worker threads have completed? What happens if a worker thread tries to access a byte of memory outside of the region's bounds? How could you avoid these problems?
  - One issue is that right now, any number of user programs can access the memory in `mymem`. Thus, while we can avoid race conditions from different threads within a user space program by using a mutex in the userspace program, to avoid clobbering between programs we would need to ensure only one program has access to the `mymem` device at a time. To do this, we could add a mutex to the open and close operations in the kernel module to ensure only one program can have an active file descriptor for the `mymem` device.
  - Right now, if you close the file descriptor before all worker threads have completed, C will happily let you keep interacting with the file, even though all the reads and writes will result in garbage values. The calls to `read/write/lseek` will all result in `-1` indicating failure, but it is up to the user code to handle these errors. An even better solution would be to refuse to compile code that allows a user to close a file descriptor and still try to use it.
  - Well there are two ways of trying to access a byte of memory. One could try to `lseek` to an illegal value, or try to `read/write` past the end of the allocated block. Both of these are safe in that the driver code will refuse these requests by returning `-EINVAL` before any damage could be done. However, it is again up to the user to notice these failures and not use the garbage values they imply.

- Data/Plot

Number of Workers vs. Percent Error (n = 150000, avg of 3 trials)

where PercentError = | (actual - expected) / expected | \* 100%



- 
- Table
- 

Number of Workers (w):	Percent Error (n = 150000, avg of 3 trials)
1	0
2	49.5
3	62.8
4	53.7
5	61.3
6	78
7	53.7
8	56.7
9	84.3
10	87.8
15	80.2
20	78.6
25	92.9
30	81
35	71
40	94.8

45	85.72
50	77.3
100	83