

GH repo: [https://github.com/evgerritz/cpsc429\\_labs/tree/2\\_p3](https://github.com/evgerritz/cpsc429_labs/tree/2_p3)

Steps to run:

On Zoo:

1. Boot server vm and forward the VNC and server ports

On Host:

1. Use VNC to connect to server
2. Enter `movenet_server/`
3. `cargo run`

On Target:

5. Connect camera to VM
6. Start yuv422 webcam with `./fake_webcam`
7. Enter `rust_movenet/`
8. `./tunnel` (to connect to the Zoo)
9. `cargo run`

## Part 3:

- 
- Acknowledgements
    - Resources provided in assignment
    - Mmap streaming example:  
<https://www.kernel.org/doc/html/v4.9/media/uapi/v4l/mmap.html>
    - Videodev2 header:  
<https://www.kernel.org/doc/html/v5.7/media/uapi/v4l/videodev.html>
    - Another v4l example:  
<https://medium.com/@athul929/capture-an-image-using-v4l2-api-5b6022d79e1d>
    - Setting up virtual yuv422 webcam:  
<https://stackoverflow.com/questions/59574987/how-to-change-mjpeg-to-yuyv422-from-a-webcam-to-a-v4l2loopback>
    - Code I translated into rust for converting from yuv422 to rgb:  
[https://github.com/kd40629rtl/yuv422\\_to\\_rgb/blob/master/yuv\\_to\\_rgb.c](https://github.com/kd40629rtl/yuv422_to_rgb/blob/master/yuv_to_rgb.c)
  - Challenges
    - Getting the virtual webcam set up
    - Creating Rust structs that were compatible with the ioctl calls was difficult. I had to use gdb to print the physical memory layout of the `v4l2_buffer` struct in order to get the alignment right.
    - Debugging failed ioctl calls
    - Converting between `libc::c_void` and `&[u8]`

- Reading a stream of RGB bytes into an OpenCV Mat
- Debugging when images were incorrectly converted
- Overview
  - The majority of the code for this assignment is in `rust_movenet/v4l_utils.rs`.
    - `v4l_utils.rs` implements the Rust interface for v4l, as will be used by the client to capture images from the webcam.
    - `main.rs` has been rewritten to use the v4l wrappers in `v4l_utils.rs`.
  - `movenet_server/main.rs` now includes a `yuv422-to-rgb24` conversion function, and the `resize_with_padding` function from `rust_movenet/util.rs`.
  - The general workflow is as follows:
    - The client uses v4l to capture a yuv422 buffer, and sends it as bytes to the server.
    - The server then converts that yuv422 buffer to rgb24, creates an OpenCV matrix, and then does the required resizing of input for the network.
    - The server then runs the DNN interpreter, and sends the results as bytes back to the client.
    - The client converts these bytes to floats, draws the output on a blank frame, and then displays the frame.